

data_jobs

March 25, 2021

1 Data Jobs

TABLE OF CONTENTS: 0. Introduction 1. Data Input 2. Data Processing/Cleaning 3. Data Visualization/EDA * 3.1. Most Frequent Tools Per Job * 3.2. Most Frequent Skills Per Job * 3.3. Job Postings by Company 4. Data Predicting/Forecasting * 4.1. Logistic Regression (One vs. All) w/ TFIDF * 4.2. Logistic Regression (Multinomial) w/ TFIDF * 4.3. Decision Tree w/ TFIDF * 4.4. K-Nearest Neighbors w/ TFIDF * 4.5. Naive Bayes w/ TFIDF * 4.6. Choose Best Model and Optimize Hyperparameters 5. Conclusion

1.1 0. Introduction:

The purpose of this project is to examine the skills and tools desired by employers for data related jobs (i.e. *Data Analyst*, *Data Scientist*, *Data Engineer*). The motivation for the project is two-fold. First, I am personally interested in data-related careers, and the skills and tools in demand from employers. Second, while job boards are helpful in searching for jobs, there is a lack of consistency in displaying which skills/tools are desired. In other words, job boards such as Indeed and LinkedIn do not have any filtering functions or ways to aggregate by skills/tools mentioned in the job advertisement. There is usually filtering functions for location, seniority, industry, etc but the filtering does not go down to the necessary level of detail for skills/tools. Furthermore, job announcements are inconsistent on where they place the text for required skills/tools. Sometimes, it is under 'Qualifications', 'Requirements', 'Skills', 'Responsibilities', or other sections. Thus, it is necessary to do some level of web scraping and text preprocessing prior to analysis.

This project will try to answer these questions: - What tools/skills are most in demand for a Data Analyst? - What tools/skills are most in demand for a Data Engineer? - What tools/skills are most in demand for a Data Scientist? - Which companies post the most data-related job openings? - Can a classifier be built which predicts job role/title (Data Analyst, Data Engineer or Data Scientist) based on job description?

1.2 1. Data Input:

Data was collected from LinkedIn and Indeed job sites via a custom, separate web scraping script.

```
[70]: # load required libraries
import os
import re
import glob
import string
```

```

import inspect
import time
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from geopy.geocoders import Nominatim
from geopy.exc import GeocoderTimedOut
from collections import Counter

# NLP libraries
import nltk
from nltk.stem import WordNetLemmatizer
from nltk.tokenize import regexp_tokenize, TweetTokenizer, sent_tokenize, ↵
    ↵word_tokenize
from nltk.tokenize.treebank import TreebankWordDetokenizer
from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer
from nltk.util import bigrams, trigrams, ngrams
from gensim.parsing.preprocessing import preprocess_documents, preprocess_string
#nltk.download('wordnet')

# sklearn libraries
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, recall_score, precision_score, ↵
    ↵f1_score, classification_report, plot_confusion_matrix, ↵
    ↵precision_recall_curve, auc, average_precision_score, ↵
    ↵plot_precision_recall_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.decomposition import PCA

```

```

[71]: def get_jobs_data(filename):

    # get current parent directory and data folder path
    par_directory = os.path.dirname(os.getcwd())
    data_directory = os.path.join(par_directory, 'data/raw')

    # retrieve job files
    files = glob.glob(os.path.join(data_directory, filename))

```

```

# create empty dataframe, loop over files and concatenate data to dataframe
df_jobs = pd.DataFrame()
for f in files:
    data = pd.read_csv(f)
    df_jobs = pd.concat([df_jobs, data], axis=0, sort=False')

# reset index
df_jobs = df_jobs.reset_index(drop=True)

return df_jobs

df_jobs = get_jobs_data('*DATA*jobs_*.csv')

# print data and length
df_jobs.head()

```

```

[71]:
      applicants
0  Be among the first 25 applicants
1  Be among the first 25 applicants
2  Be among the first 25 applicants
3      114 applicants      Prescriptive Health, Inc.
4      30 applicants      Source One Technical Solutions

      company \
0  Be among the first 25 applicants      Microf
1  Be among the first 25 applicants      NaN
2  Be among the first 25 applicants      Adobe
3      114 applicants      Prescriptive Health, Inc.
4      30 applicants      Source One Technical Solutions

      company_rating date      date_posted      date_scraped employment_type \
0      NaN      NaN      7 hours ago      2021-03-18 08:10:59      Full-time
1      NaN      NaN      10 hours ago      2021-03-18 08:10:59      Internship
2      NaN      NaN      2 hours ago      2021-03-18 08:10:59      Full-time
3      NaN      NaN      14 hours ago      2021-03-18 08:10:59      Full-time
4      NaN      NaN      12 hours ago      2021-03-18 08:10:59      Contract

      industries      job_function \
0      Marketing and Advertising      Information Technology
1      NaN      Information Technology
2      Marketing and Advertising      Information Technology
3      Information Technology and Services      Information Technology
4      Information Technology and Services      Information Technology

      job_text \
0      ['Eligible for participation in the Compan...
1      ['Qualifications', 'Understand the day-to-...
2      ['Define, measure and track key metrics to...
3      ['\tPrescriptive Health is putting an end ...
4      ['Qualifications', '• Analytical and data ...

      job_title      location \
0      Data Analyst      Roswell, GA

```

```

1          Data Analyst  New York City Metropolitan Area
2  2021 Intern - Data Analyst      California, United States
3          Data Analyst      Redmond, WA
4          Data Analyst      Richboro, PA

```

```

    seniority_level
0  Mid-Senior level
1    Entry level
2    Internship
3  Mid-Senior level
4    Associate

```

1.3 2. Data Processing/Cleaning:

```

[72]: # view descriptive info on dataframe
df_jobs.info()
df_jobs.describe()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148882 entries, 0 to 148881
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   applicants             66417 non-null  object
1   company                148341 non-null object
2   company_rating         51562 non-null  object
3   date                   80727 non-null  object
4   date_posted            65432 non-null  object
5   date_scraped           68155 non-null  object
6   employment_type        68155 non-null  object
7   industries             67733 non-null  object
8   job_function           68034 non-null  object
9   job_text               148882 non-null object
10  job_title              142640 non-null object
11  location                132605 non-null object
12  seniority_level        68155 non-null  object
dtypes: object(13)
memory usage: 14.8+ MB

```

```

[72]:
count      applicants  company  company_rating  \
unique              178    19535             5481
top    Be among the first 25 applicants  Facebook  3.9 out of 5
freq              50729     2130             3330

count      date  date_posted      date_scraped  employment_type  \
count      80727     65432             68155             68155

```

unique	5316	103	114	7
top	2020-07-27 12:04:05	5 hours ago	2020-09-03 10:26:32	Full-time
freq	19	3575	1000	61401

	industries	job_function	job_text \
count	67733	68034	148882
unique	137	84	79058
top	Information Technology and Services	Engineering	[]
freq	22862	25019	5245

	job_title	location	seniority_level
count	142640	132605	68155
unique	29997	5951	7
top	Data Scientist	New York, NY	Mid-Senior level
freq	14214	6165	20342

```
[73]: # clean jobs description data
def clean_jobs(df):

    # clean job text data with empty list, reg expressions, and appending to
    ↪ list
    clean_text = []

    for x in df['job_text']:

        x = re.sub(r'(?<=[.,])(?=[^\s])', r' ', x)
        x = re.sub(r'<[A-Za-z/*\>+', '', x)
        x = re.sub(r'\\xa0', '', x)
        x = re.sub(r'\\n', '', x)
        x = re.sub(r'Data Analyst', '', x)
        x = re.sub(r'Data Engineer', '', x)
        x = re.sub(r'Data Scientist', '', x)

        clean_text.append(x)

    df['clean_text'] = clean_text

    # clean date columns by filling in missing values, converting to pd
    ↪ datetime format
    df['date'].update(df.pop('date_scraped'))
    df['date'] = pd.to_datetime(df['date'])

    # clean rating columns by extracting rating string
    df['company_rating'] = [x.split('out')[0] if type(x) != float else x for x
    ↪ in df['company_rating']]

    # drop NA, duplicates, and unnecessary columns
```

```

df = df.drop_duplicates(subset=['clean_text', 'job_title'])
df = df.dropna(subset=['clean_text', 'job_title'])
df = df[df['clean_text'] != '[]']
df = df.drop(columns=['date_posted', 'seniority_level', 'applicants', '
→ 'job_function', 'employment_type'])

# filter for jobs with description length greater than 10 words
df['job_text_length'] = df['clean_text'].apply(lambda x: len(x))
df = df[df['job_text_length'] >= 10]

# reset dataframe index
df.reset_index(drop=True, inplace=True)

return df

df_jobs = clean_jobs(df_jobs)

# print number of jobs and data sample
print('Number of Jobs: {}'.format(len(df_jobs)))
df_jobs.head()

```

Number of Jobs: 77694

```

[73]:

```

	company	company_rating	date	\
0	Microf	NaN	2021-03-18 08:10:59	
1	NaN	NaN	2021-03-18 08:10:59	
2	Adobe	NaN	2021-03-18 08:10:59	
3	Prescriptive Health, Inc.	NaN	2021-03-18 08:10:59	
4	Source One Technical Solutions	NaN	2021-03-18 08:10:59	

	industries	\
0	Marketing and Advertising	
1	NaN	
2	Marketing and Advertising	
3	Information Technology and Services	
4	Information Technology and Services	

	job_text	\
0	['', 'Eligible for participation in the Compan...	
1	['', 'Qualifications', 'Understand the day-to-...	
2	['', 'Define, measure and track key metrics to...	
3	['', '\tPrescriptive Health is putting an end ...	
4	['', 'Qualifications', '• Analytical and data ...	

	job_title	location	\
0	Data Analyst	Roswell, GA	
1	Data Analyst	New York City Metropolitan Area	

2	2021 Intern - Data Analyst	California, United States
3	Data Analyst	Redmond, WA
4	Data Analyst	Richboro, PA

		clean_text	job_text_length
0	['', 'Eligible for participation in the Compan...		6529
1	['', 'Qualifications', 'Understand the day-to-...		1169
2	['', 'Define, measure and track key metrics to...		1940
3	['', '\tPrescriptive Health is putting an end ...		3053
4	['', 'Qualifications', '• Analytical and data ...		1100

```
[74]: # function used to geocode locations and override timeout error
#def geocode_location(location):
#    time.sleep(1)
#    geopy = Nominatim(user_agent="my_project")
#    try:
#        return geopy.geocode(location,exactly_one=True, country_codes='us')
#    except GeocoderTimedOut:
#        return do_geocode(location)

#df_jobs['geocoded_location']=df_jobs['location'].apply(lambda x:
#    ↪geocode_location(x) if x != None else None)

# create latitude and longitude column from geocoded location
#df_jobs['latitude'] = df_jobs['geocoded_location'].apply(lambda x: x[1][0] if
#    ↪x != None else None)
#df_jobs['longitude'] = df_jobs['geocoded_location'].apply(lambda x: x[1][1] if
#    ↪x != None else None)
#print(df_jobs.head())
```

```
[75]: # filter for job titles with Data Scientist, Data Engineer, or Data Analyst
def filter_jobs(df):

    # filter for data scientist, data engineer, or data analyst
    df_jobs = df[(df['job_title'].str.contains('Data Scientist', case=False)) |
    ↪(df['job_title'].str.contains('Data Engineer', case=False)) |
    ↪(df['job_title'].str.contains('Data Analyst', case=False))].copy()

    # add identifying column for data scientist (#0) and data engineer (#1), or
    ↪data analyst (#2)
    df_jobs['label'] = df_jobs['job_title'].apply(lambda x: 0 if 'Scientist' in
    ↪x or 'scientist' in x or 'SCIENTIST' in x else 1 if 'Engineer' in x or
    ↪'engineer' in x or 'ENGINEER' in x else 2 if 'Analyst' in x or 'analyst' in
    ↪x or 'ANALYST' in x else '')
    df_jobs = df_jobs.dropna(subset=['label'])
    df_jobs.reset_index(inplace=True, drop=True)
```

```

    # print number of jobs and counts of each job title
    print('\nCounts of Job Titles (0=Data Scientist, 1=Data Engineer, 2=Data_
→Analyst): \n\n{}'.format(df_jobs['label'].value_counts()))

    return df_jobs

df_jobs = filter_jobs(df_jobs)

```

Counts of Job Titles (0=Data Scientist, 1=Data Engineer, 2=Data Analyst):

```

1    15206
0    14692
2     7364
Name: label, dtype: int64

```

1.4 3. Data Visualization/EDA

1.4.1 3.1 Most Frequent Tools per Job

```

[76]: # visualize job skills/tools per job title
def data_tools(df, title):

    # filter jobs for job title & count
    #df_jobs = df[df['label'] == label]
    df_jobs = df[(df['job_title'].str.contains(title, case=False))]
    num_jobs = len(df_jobs)

    # Tokenize the article: tokens
    tokens = [word_tokenize(x) for x in df_jobs['job_text']]
    tokens = [item for sublist in tokens for item in sublist]

    # Convert the tokens into lowercase: lower_tokens
    lower_tokens = [t.lower() for t in tokens]

    # # Retain alphabetic words: alpha_only
    alpha_only = [t for t in lower_tokens if t.isalpha()]

    # set stop words
    stop_words = set(stopwords.words('english'))

    # # Remove all stop words: no_stops
    no_stops = [t for t in alpha_only if t not in stop_words]

    # # Instantiate the WordNetLemmatizer
    wordnet_lemmatizer = WordNetLemmatizer()

```



```

# # Lemmatize all tokens into a new list: lemmatized
lemmatized = [wordnet_lemmatizer.lemmatize(t) for t in no_stops]
print(type(lemmatized))

# # Create the bag-of-words: bow
bow = Counter(lemmatized)

# create dataframe from dictionary
df_count = pd.DataFrame.from_dict(bow, orient='index').reset_index()
df_count.columns = ['keywords', 'counts']
df_count = df_count.sort_values(by='counts', ascending=False,
→ignore_index=True)
df_count['avg_frequency_per_job'] = df_count['counts'] / num_jobs

# create list of data tools
data_tools = ['airflow', 'azure', 'aws', 'bi', 'bigquery', 'c', 'c++',
→'d3', 'docker', 'ec2', 'excel', 'git', 'hadoop', 'hive',
        'java', 'javascript', 'jenkins', 'jupyter', 'kafka', 'keras',
→'kubernetes', 'linux', 'luigi', 'matlab', 'mongodb',
        'perl', 'python', 'pytorch', 'r', 'react', 'redshift',
→'ruby', 'sas', 'scala', 'scikit-learn', 'sql', 'spark', 'tableau',
→'tensorflow']

# filter keyword counts for data tools
df_count = df_count[df_count['keywords'].isin(data_tools)]
df_count.reset_index(drop=True, inplace=True)

# Print the 20 most common tools
print('\n' + title + ' Top Keywords:\n\n', df_count.iloc[:20])

# plot the 20 most common tools
sns.set()
fig, ax = plt.subplots(dpi=300) #, figsize = (3, 5), )
sns.barplot(x="avg_frequency_per_job", y="keywords", data=df_count.iloc[:
→20], palette="Blues_d")
plt.title(label = "" + title + "" + ' Keywords on LinkedIn/Indeed',
→fontsize=13)
plt.xlabel('Frequency (per job posting)', fontsize=8)
plt.ylabel('Keywords', fontsize=8)

# ax.set(title = "" + title + "" + ' Keywords on LinkedIn/Indeed')
plt.show()

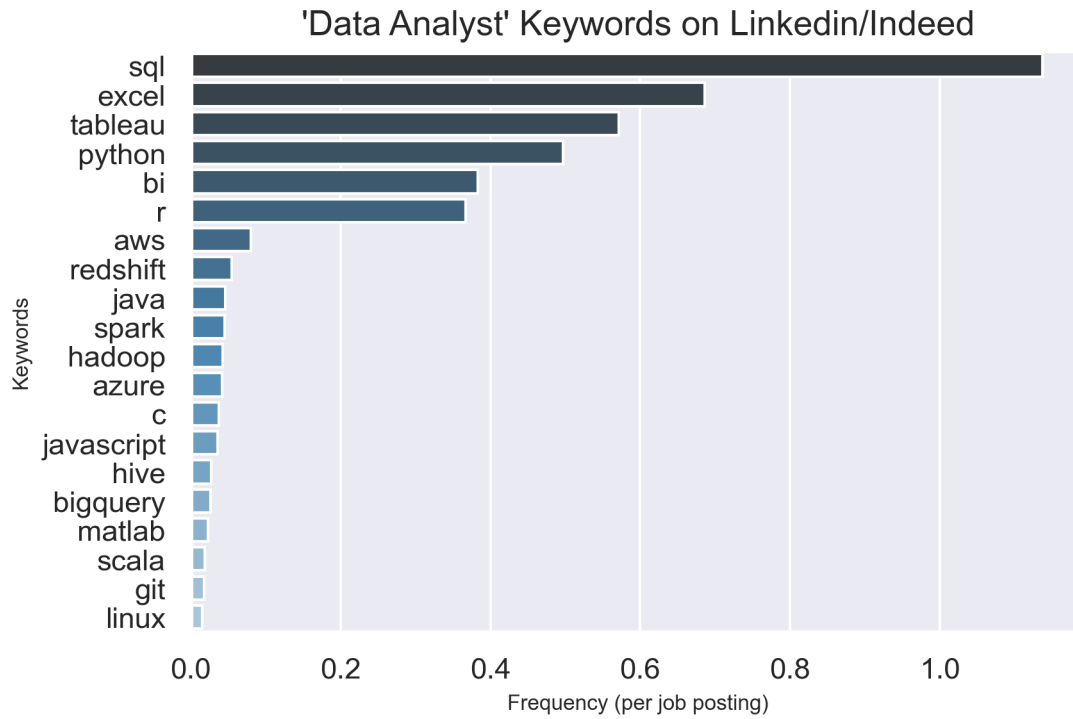
```

```
[77]: data_tools(df_jobs, 'Data Analyst')
```

```
<class 'list'>
```

Data Analyst Top Keywords:

	keywords	counts	avg_frequency_per_job
0	sql	8649	1.137428
1	excel	5215	0.685823
2	tableau	4343	0.571147
3	python	3779	0.496975
4	bi	2915	0.383351
5	r	2787	0.366518
6	aws	608	0.079958
7	redshift	416	0.054708
8	java	346	0.045502
9	spark	343	0.045108
10	hadoop	322	0.042346
11	azure	313	0.041163
12	c	286	0.037612
13	javascript	272	0.035771
14	hive	203	0.026696
15	bigquery	198	0.026039
16	matlab	174	0.022883
17	scala	138	0.018148
18	git	136	0.017885
19	linux	111	0.014598




```
[78]: data_tools(df_jobs, 'Data Engineer')
```

```
<class 'list'>
```

Data Engineer Top Keywords:

	keywords	counts	avg_frequency_per_job
0	sql	23581	1.551586
1	python	17491	1.150875
2	aws	14321	0.942295
3	spark	11438	0.752599
4	azure	9853	0.648309
5	java	6980	0.459271
6	hadoop	6704	0.441111
7	bi	5089	0.334847
8	redshift	4803	0.316028
9	kafka	4765	0.313528
10	scala	4614	0.303593
11	tableau	3911	0.257336
12	hive	3823	0.251546
13	airflow	3656	0.240558
14	r	2932	0.192920
15	git	1999	0.131530
16	docker	1886	0.124095
17	linux	1712	0.112646
18	c	1645	0.108238
19	mongodb	1357	0.089288




```
[79]: data_tools(df_jobs, 'Data Scientist')
```

```
<class 'list'>
```

Data Scientist Top Keywords:

	keywords	counts	avg_frequency_per_job
0	python	18520	1.268146
1	r	12715	0.870652
2	sql	12551	0.859422
3	spark	4626	0.316763
4	tableau	4452	0.304848
5	aws	4036	0.276363
6	hadoop	3356	0.229800
7	java	3255	0.222884
8	tensorflow	2566	0.175705
9	scala	2060	0.141057
10	hive	1804	0.123528
11	azure	1763	0.120720
12	excel	1761	0.120583
13	matlab	1735	0.118803
14	bi	1720	0.117776
15	pytorch	1569	0.107436
16	c	1245	0.085251

17	git	1059	0.072514
18	linux	865	0.059230
19	docker	760	0.052041



1.4.2 3.2 Most Frequent Skills per Job

```
[80]: # preprocess job description string
def preprocess_text(df):

    df = df.copy()

    processed_text = []

    for x in df['clean_text']:

        ## Convert the tokens into lowercase: lower_tokens
        lower_tokens = [x.lower()]

        ## Convert tokens
        tokenize = [word_tokenize(x) for x in lower_tokens]

        ### Retain alphabetic words: alpha_only
        alpha_only = [i for item in tokenize for i in item if i.isalpha()]
```

```

# set stop words
stop_words = set(stopwords.words('english'))

# # Remove all stop words: no_stops
no_stops = [i for i in alpha_only if i not in stop_words]

# # Instantiate the WordNetLemmatizer
wordnet_lemmatizer = WordNetLemmatizer()

# # Lemmatize all tokens into a new list: lemmatized
lemmatized = [wordnet_lemmatizer.lemmatize(i) for i in no_stops]

processed_text.append(lemmatized)

df['processed_text'] = processed_text

return df

df_jobs = preprocess_text(df_jobs)

df_jobs.head()

```

```

[80]:
      company company_rating      date \
0      Microf              NaN 2021-03-18 08:10:59
1              NaN              NaN 2021-03-18 08:10:59
2      Adobe              NaN 2021-03-18 08:10:59
3  Prescriptive Health, Inc.      NaN 2021-03-18 08:10:59
4  Source One Technical Solutions      NaN 2021-03-18 08:10:59

      industries \
0  Marketing and Advertising
1              NaN
2  Marketing and Advertising
3  Information Technology and Services
4  Information Technology and Services

      job_text \
0  ['', 'Eligible for participation in the Compan...
1  ['', 'Qualifications', 'Understand the day-to-...
2  ['', 'Define, measure and track key metrics to...
3  ['', '\tPrescriptive Health is putting an end ...
4  ['', 'Qualifications', '• Analytical and data ...

      job_title      location \
0  Data Analyst      Roswell, GA
1  Data Analyst  New York City Metropolitan Area

```

2	2021 Intern - Data Analyst	California, United States
3	Data Analyst	Redmond, WA
4	Data Analyst	Richboro, PA

	clean_text	job_text_length	label	\
0	['', 'Eligible for participation in the Compan...	6529	2	
1	['', 'Qualifications', 'Understand the day-to-...	1169	2	
2	['', 'Define, measure and track key metrics to...	1940	2	
3	['', '\tPrescriptive Health is putting an end ...	3053	2	
4	['', 'Qualifications', '• Analytical and data ...	1100	2	

	processed_text
0	[participation, company, group, medical, plan,...
1	[understand, issue, business, face, better, un...
2	[measure, track, key, metric, guide, execution...
3	[health, putting, end, nation, prescription, d...
4	[analytical, data, visualization, skill, requi...

```
[81]: def ngram_generator(df, n, title):

    # filter for job title
    df_jobs = df[(df['job_title'].str.contains(title, case=False))]

    # create empty list
    ngram_list = []

    # loop over every row in df_jobs['clean_text']
    for x in df_jobs['processed_text']:

        # join each list of strings into sentence
        joined = [' '.join(x)]

        # create list of trigrams within each row
        ngram = [list(ngrams(item.split(), n)) for item in joined]

        # append list of trigrams to empty list
        ngram_list.append(ngram)

    # extract item within each sublist
    ngram_list = [item for sublist in ngram_list for item in sublist]

    # extract each item within sublist again for Counter (list is unhashable)
    ngram_list = [item for sublist in ngram_list for item in sublist]

    # print the 20 most common grams via Counter function

    return list(Counter(ngram_list).most_common(20))
```

1.4.3 Data Analyst

```
[82]: ngram_generator(df_jobs, 2, 'Data Analyst')
```

```
[82]: [ (('indeed', 'center'), 4958),  
      (('data', 'analysis'), 4766),  
      (('year', 'experience'), 4627),  
      (('job', 'review'), 3781),  
      (('review', 'salary'), 3781),  
      (('communication', 'skill'), 3470),  
      (('data', 'set'), 2654),  
      (('data', 'visualization'), 2597),  
      (('personal', 'information'), 2520),  
      (('computer', 'science'), 2503),  
      (('job', 'company'), 2481),  
      (('company', 'certification'), 2480),  
      (('salary', 'resume'), 2479),  
      (('lab', 'advice'), 2479),  
      (('advice', 'job'), 2479),  
      (('certification', 'event'), 2479),  
      (('event', 'indeed'), 2479),  
      (('center', 'indeed'), 2479),  
      (('indeed', 'sell'), 2479),  
      (('sell', 'personal'), 2479)]
```

1.4.4 Data Engineer

```
[83]: ngram_generator(df_jobs, 2, 'Data Engineer')
```

```
[83]: [ (('year', 'experience'), 17273),  
      (('data', 'pipeline'), 13777),  
      (('big', 'data'), 10145),  
      (('computer', 'science'), 9833),  
      (('data', 'warehouse'), 8137),  
      (('experience', 'working'), 6087),  
      (('data', 'engineering'), 5705),  
      (('best', 'practice'), 5316),  
      (('indeed', 'center'), 5290),  
      (('machine', 'learning'), 5190),  
      (('data', 'modeling'), 5123),  
      (('data', 'quality'), 5099),  
      (('data', 'set'), 5083),  
      (('degree', 'computer'), 5077),  
      (('experience', 'data'), 4989),  
      (('data', 'source'), 4939),  
      (('communication', 'skill'), 4908),  
      (('data', 'science'), 4892),
```



```
((('data', 'model'), 4689),
 (('job', 'review'), 4160])
```

1.4.5 Data Scientist

```
[84]: ngram_generator(df_jobs, 2, 'Data Scientist')
```

```
[84]: [((('machine', 'learning'), 23459),
 (('data', 'science'), 20154),
 (('year', 'experience'), 11336),
 (('computer', 'science'), 10376),
 (('data', 'analysis'), 6340),
 (('data', 'set'), 5859),
 (('communication', 'skill'), 5749),
 (('indeed', 'center'), 5020),
 (('data', 'visualization'), 4594),
 (('experience', 'data'), 4501),
 (('job', 'review'), 4374),
 (('review', 'salary'), 4374),
 (('related', 'field'), 4358),
 (('big', 'data'), 4308),
 (('python', 'r'), 4240),
 (('deep', 'learning'), 4175),
 (('r', 'python'), 3815),
 (('data', 'mining'), 3754),
 (('data', 'source'), 3749),
 (('programming', 'language'), 3533)]
```

1.4.6 3.3 Most Frequent Companies Posting Jobs

```
[85]: def job_companies(df, title):

    # replace various Amazon company names with Amazon
    df['company'] = df['company'].replace(to_replace = 'Amazon Web Services,
    ↳(AWS)', value = 'Amazon')
    df['company'] = df['company'].replace(to_replace = 'Amazon.com Services,
    ↳LLC', value = 'Amazon')
    df['company'] = df['company'].replace(to_replace = 'Amazon Web Services,
    ↳Inc.', value = 'Amazon')

    # print value counts for companies to determine top 20 companies for each
    ↳type of job posting
    title_companies = df[(df['job_title'].str.contains(title,
    ↳case=False))]['company'].value_counts().iloc[:20]
    print('\nMost Frequent Companies for {} Job Postings: \n\n{}'.format(title,
    ↳title_companies))
```

```

# plot most frequent companies
sns.set()
fig, ax = plt.subplots(figsize = (12,8), dpi=250)
sns.barplot(x = title_companies.iloc[:10].index, y= title_companies.iloc[:
↪10].values, palette="Blues_d")
plt.title(label = "Most Frequent Companies for " + title + " Jobs",
↪fontsize=14)
plt.xlabel('Companies', fontsize=14)
plt.ylabel('No. of Postings', fontsize=14)
plt.xticks(rotation=45, fontsize=12)
plt.show()

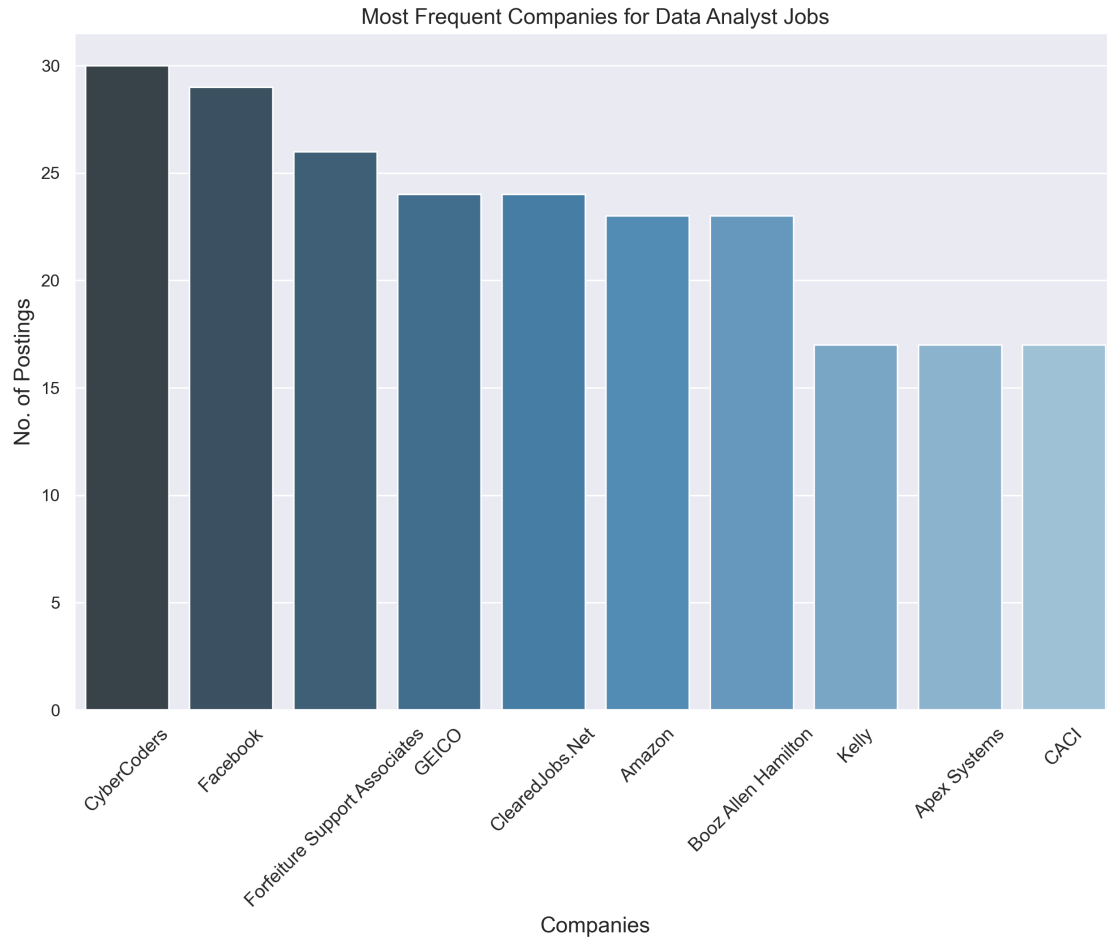
```

1.4.7 Data Analyst

```
[86]: job_companies(df_jobs, 'Data Analyst')
```

Most Frequent Companies for Data Analyst Job Postings:

CyberCoders	30
Facebook	29
Forfeiture Support Associates	26
GEICO	24
ClearedJobs.Net	24
Amazon	23
Booz Allen Hamilton	23
Kelly	17
Apex Systems	17
CACI	17
Intuit	17
Insight Global	16
Tesla	15
Robert Half	14
Kforce Inc	14
Piper Companies	14
Guidehouse	13
UnitedHealth Group	13
Vanguard	13
Microsoft	13
Name: company, dtype: int64	



1.4.8 Data Engineer

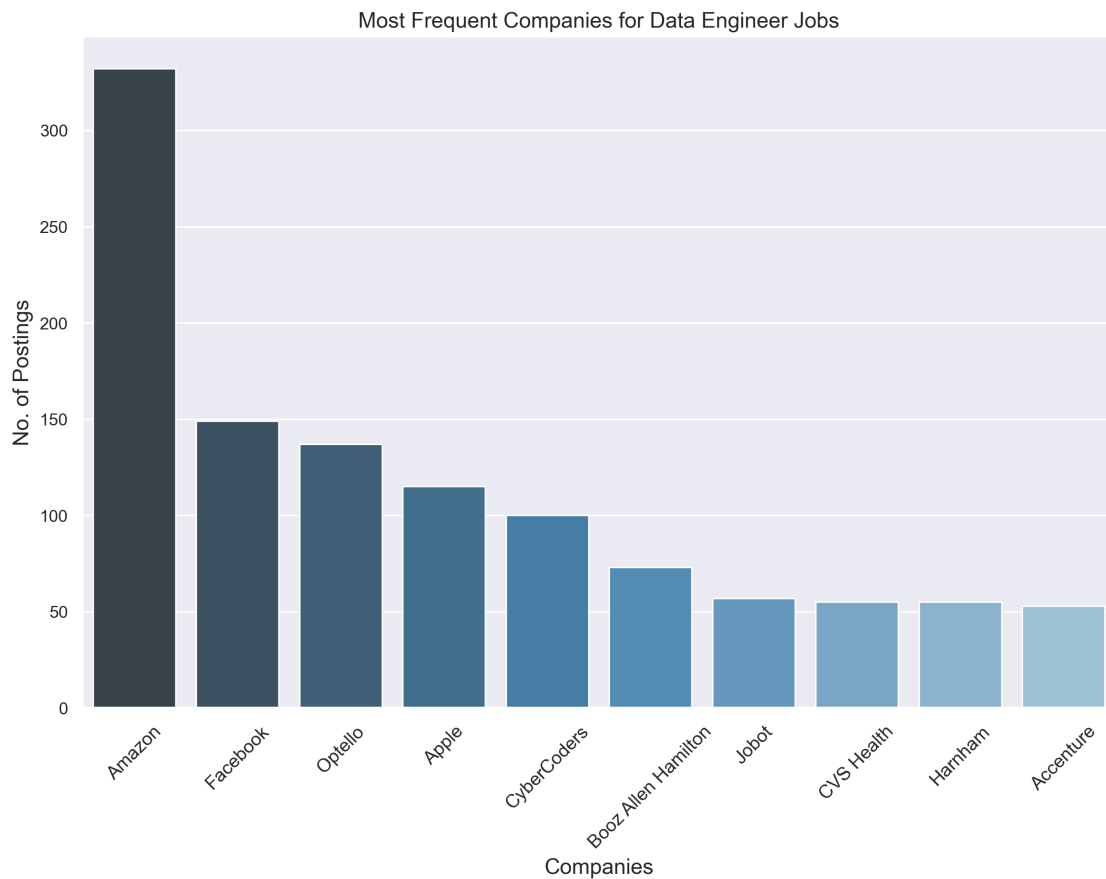
```
[87]: job_companies(df_jobs, 'Data Engineer')
```

Most Frequent Companies for Data Engineer Job Postings:

Amazon	332
Facebook	149
Optello	137
Apple	115
CyberCoders	100
Booz Allen Hamilton	73
Jobot	57
CVS Health	55
Harnham	55
Accenture	53
Cognizant	51

Deloitte	47
USAA	47
Quantitative Systems	46
Tesla	45
UnitedHealth Group	45
Apex Systems	44
Idexcel	44
Brooksource	39
Capgemini	38

Name: company, dtype: int64



1.4.9 Data Scientist

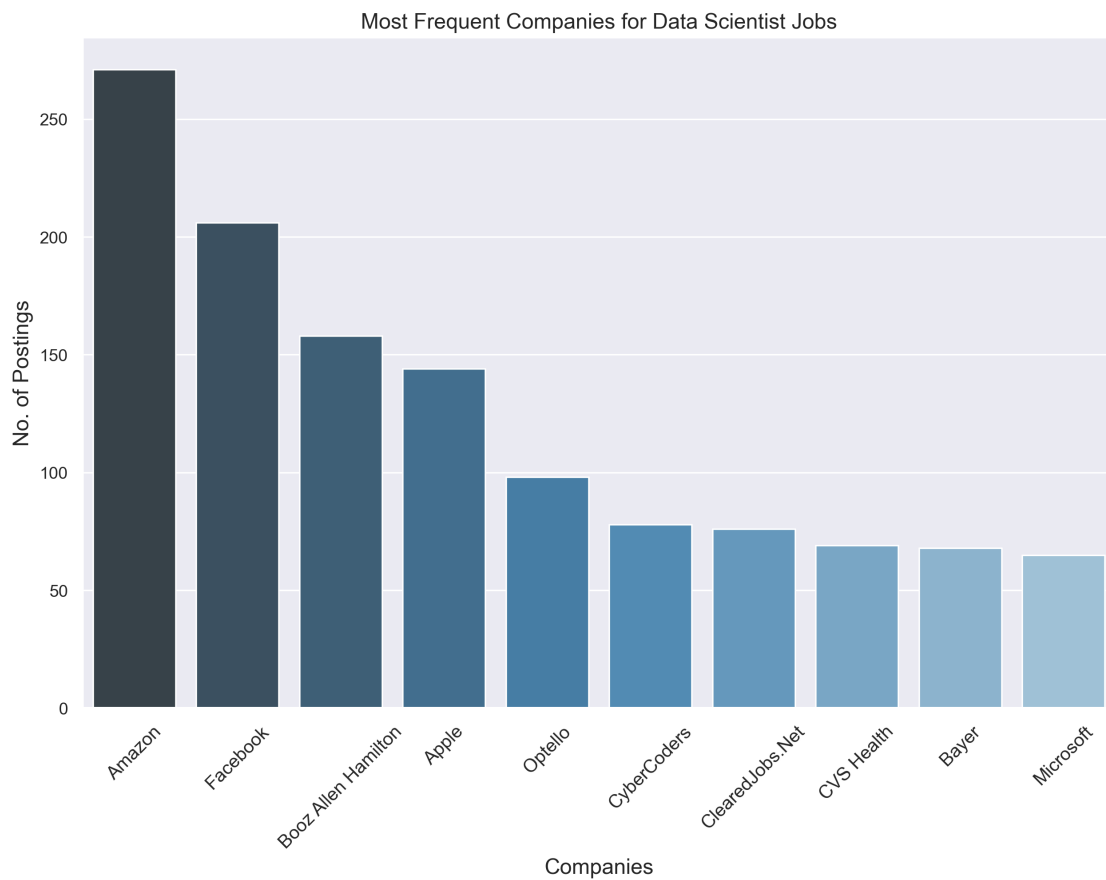
```
[88]: job_companies(df_jobs, 'Data Scientist')
```

Most Frequent Companies for Data Scientist Job Postings:

Amazon	271
Facebook	206

Booz Allen Hamilton	158
Apple	144
Optello	98
CyberCoders	78
ClearedJobs.Net	76
CVS Health	69
Bayer	68
Microsoft	65
IBM	65
Leidos	60
Deloitte	57
JPMorgan Chase Bank, N.A.	53
Harnham	53
Guidehouse	52
Walmart	51
Spotify	50
Mitre Corporation	49
PayPal	48

Name: company, dtype: int64



1.5 4. Data Prediction

```
[89]: # # make series of job text for just Data Engineer and Data Scientist roles
def remove_obvious_words(df):

    updated_words = []

    for x in df['processed_text']:

        x = str(x)
        x = x.replace('scientist', '')
        x = x.replace('engineer', '')
        x = x.replace('analyst', '')

        # append updated strings to list
        updated_words.append(x)

    df['processed_text'] = updated_words

    return df

df_jobs = remove_obvious_words(df_jobs)
df_jobs.head()
```

```
[89]:
```

	company	company_rating	date \
0	Microf	NaN	2021-03-18 08:10:59
1	NaN	NaN	2021-03-18 08:10:59
2	Adobe	NaN	2021-03-18 08:10:59
3	Prescriptive Health, Inc.	NaN	2021-03-18 08:10:59
4	Source One Technical Solutions	NaN	2021-03-18 08:10:59

	industries \
0	Marketing and Advertising
1	NaN
2	Marketing and Advertising
3	Information Technology and Services
4	Information Technology and Services

	job_text \
0	['', 'Eligible for participation in the Compan...
1	['', 'Qualifications', 'Understand the day-to-...
2	['', 'Define, measure and track key metrics to...
3	['', '\tPrescriptive Health is putting an end ...
4	['', 'Qualifications', '• Analytical and data ...

	job_title	location \
0	Data Analyst	Roswell, GA

```

1          Data Analyst   New York City Metropolitan Area
2 2021 Intern - Data Analyst   California, United States
3          Data Analyst   Redmond, WA
4          Data Analyst   Richboro, PA

```

```

              clean_text  job_text_length  label  \
0  [' ', 'Eligible for participation in the Compan...      6529      2
1  [' ', 'Qualifications', 'Understand the day-to-...      1169      2
2  [' ', 'Define, measure and track key metrics to...      1940      2
3  [' ', '\tPrescriptive Health is putting an end ...      3053      2
4  [' ', 'Qualifications', '• Analytical and data ...      1100      2

```

```

              processed_text
0  ['participation', 'company', 'group', 'medical...
1  ['understand', 'issue', 'business', 'face', 'b...
2  ['measure', 'track', 'key', 'metric', 'guide',...
3  ['health', 'putting', 'end', 'nation', 'prescr...
4  ['analytical', 'data', 'visualization', 'skill...

```

```

[90]: # create feature and target series from dataframe
df_feature = df_jobs[(df_jobs['label'] == 0) | (df_jobs['label'] == 1) |
    ↳(df_jobs['label'] == 2)].loc[:, 'processed_text']
df_target = df_jobs[(df_jobs['label'] == 0) | (df_jobs['label'] == 1) |
    ↳(df_jobs['label'] == 2)].loc[:, 'label']

```

```

[91]: # apply TF-IDF based feature representation
tfidf_vectorizer = TfidfVectorizer()
df_feature_TFIDF = tfidf_vectorizer.fit_transform(df_feature)

# split train/test data 80/20
X_train, X_test, y_train, y_test = train_test_split(df_feature_TFIDF,
    df_target,
    train_size=0.8,
    random_state=20)

```

```

[92]: # create empty dataframe to store model results and scores
model_results = pd.DataFrame(columns=['model', 'accuracy', 'precision',
    ↳'recall', 'f1'])
model_results

```

```

[92]: Empty DataFrame
Columns: [model, accuracy, precision, recall, f1]
Index: []

```

1.5.1 4.1 Logistic Regression (One vs. All) w/ TFIDF (Model 1)

```
[93]: # define logistic regression, fit to model
logreg_OVR_TFIDF = LogisticRegression(multi_class = 'ovr')
#log_reg_CV = LogisticRegression(solver='liblinear', penalty='l1')
logreg_OVR_TFIDF.fit(X_train, y_train)
```

```
[93]: LogisticRegression(multi_class='ovr')
```

```
[94]: # compute y-prediction and accuracy, recall, precision, and f1 scores
y_pred = logreg_OVR_TFIDF.predict(X_test)
mod1_acc = accuracy_score(y_test, y_pred)
mod1_recall = recall_score(y_test, y_pred, average='macro')
mod1_precision = precision_score(y_test, y_pred, average='macro')
mod1_f1 = f1_score(y_test, y_pred, average='macro')

# print accuracy, recall, precision, f1 score, detailed report
print('-----')
print('Model 1 accuracy: ', round(mod1_acc,3))
print('Model 1 recall: ', round(mod1_recall,3))
print('Model 1 precision: ', round(mod1_precision,3))
print('Model 1 f1 : ', round(mod1_f1,3))
print('Model 1 classification report: \n',classification_report(y_test, y_pred))
```

```
-----
Model 1 accuracy:  0.872
Model 1 recall:    0.853
Model 1 precision: 0.863
Model 1 f1 :      0.857
Model 1 classification report:
              precision    recall  f1-score   support

     0           0.87         0.89         0.88         2900
     1           0.90         0.91         0.90         3025
     2           0.82         0.76         0.79         1528

 accuracy                   0.87         7453
 macro avg           0.86         0.85         0.86         7453
weighted avg           0.87         0.87         0.87         7453
```

```
[95]: # append model results to dataframe
      # append to dataframe
model_results = model_results.append({'model': 'LogReg_OVR_TFIDF',
                                     'accuracy':round(mod1_acc,3),
                                     'recall': round(mod1_recall,3),
                                     'precision': round(mod1_precision,3),
```



```

        'f1': round(mod1_f1,3)}, ignore_index=True)

model_results

```

```

[95]:
      model  accuracy  precision  recall    f1
0  LogReg_OVR_TFIDF    0.872    0.863   0.853  0.857

```

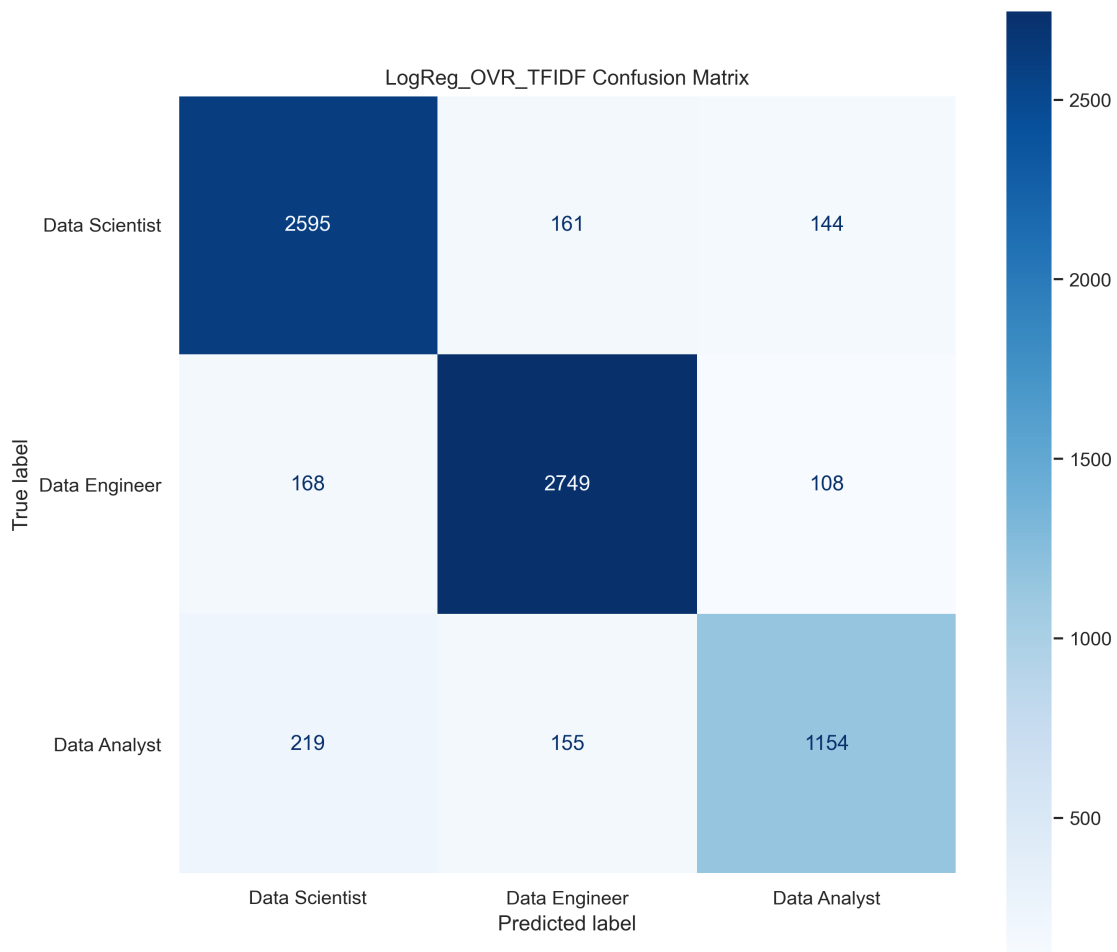
```

[96]: fig, ax = plt.subplots(figsize=(10, 10), dpi=300)

plot_confusion_matrix(logreg_OVR_TFIDF, X_test, y_test,
                      display_labels=['Data Scientist', 'Data Engineer', 'Data Analyst'],
                      cmap=plt.cm.Blues, ax=ax)
ax.set_title('LogReg_OVR_TFIDF Confusion Matrix')

plt.grid(None)
plt.show()

```



1.5.2 4.2 Logistic Regression (Multinomial) w/ TF-IDF (Model 2)

```
[97]: # define logistic regression, fit to model
```

```
logreg_multi_TFIDF = LogisticRegression(multi_class = 'multinomial')
logreg_multi_TFIDF.fit(X_train, y_train)
```

```
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[97]: LogisticRegression(multi_class='multinomial')
```

```
[98]: # compute y-prediction and accuracy, recall, precision, and f1 scores
```

```
y_pred = logreg_multi_TFIDF.predict(X_test)
mod2_acc = accuracy_score(y_test, y_pred)
mod2_recall = recall_score(y_test, y_pred, average = 'macro')
mod2_precision = precision_score(y_test, y_pred, average = 'macro')
mod2_f1 = f1_score(y_test, y_pred, average = 'macro')

# print accuracy, recall, precision, f1 score, detailed report
print('-----')
print('Model 2 accuracy: ', round(mod2_acc,3))
print('Model 2 recall: ', round(mod2_recall,3))
print('Model 2 precision: ', round(mod2_precision,3))
print('Model 2 f1 : ', round(mod2_f1,3))
print('Model 2 classification report: \n',classification_report(y_test, y_pred))
```

```
-----
Model 2 accuracy:  0.88
```

```
Model 2 recall:   0.865
```

```
Model 2 precision: 0.87
```

```
Model 2 f1 :     0.867
```

```
Model 2 classification report:
```

	precision	recall	f1-score	support
0	0.88	0.90	0.89	2900
1	0.91	0.91	0.91	3025
2	0.82	0.78	0.80	1528

accuracy			0.88	7453
macro avg	0.87	0.86	0.87	7453
weighted avg	0.88	0.88	0.88	7453

```
[99]: # append model results to dataframe
model_results = model_results.append({'model': 'LogReg_Multi_TFIDF',
                                     'accuracy': round(mod2_acc,3),
                                     'recall': round(mod2_recall,3),
                                     'precision': round(mod2_precision,3),
                                     'f1': round(mod2_f1,3)}, ignore_index=True)

model_results
```

```
[99]:
```

	model	accuracy	precision	recall	f1
0	LogReg_OVR_TFIDF	0.872	0.863	0.853	0.857
1	LogReg_Multi_TFIDF	0.880	0.870	0.865	0.867

```
[100]: fig, ax = plt.subplots(figsize=(10, 10), dpi=300)

plot_confusion_matrix(logreg_multi_TFIDF, X_test, y_test,
                      display_labels=['Data Scientist', 'Data Engineer', 'Data Analyst'],
                      cmap=plt.cm.Blues, ax=ax)
ax.set_title('LogReg_Multi_TFIDF Confusion Matrix')

plt.grid(None)
plt.show()
```



```
[101]: # get importance
importance = logreg_multi_TFIDF.coef_[0]

# create empty dataframe
feature_list = []

# summarize feature importance
for i,v in enumerate(importance):
    if v > 0.5:
        feature_list.append(dict({'number': i, 'feature': tfidf_vectorizer.
→get_feature_names()[i], 'weight': v}))

    if v < -0.5:
        feature_list.append(dict({'number': i, 'feature': tfidf_vectorizer.
→get_feature_names()[i], 'weight': v}))
```

```
df_important_features = pd.DataFrame(feature_list, columns=['number',
↳ 'feature', 'weight'])
print('Top Words (Features) for Predicting Data Scientist job: \n\n',
↳ df_important_features.sort_values(by='weight', ascending=False).iloc[:10])
print('\nTop Words (Features) for Predicting Data Engineering job: \n\n',
↳ df_important_features.sort_values(by='weight', ascending=True).iloc[:10])
```

Top Words (Features) for Predicting Data Scientist job:

	number	feature	weight
907	48453	science	7.777486
584	29640	learning	4.649844
656	33423	model	4.069614
979	52744	statistic	3.172504
741	39188	phd	3.087836
384	19632	experiment	3.069617
816	43674	python	3.014489
603	30837	machine	2.977213
34	1472	algorithm	2.881187
980	52747	statistical	2.796520

Top Words (Features) for Predicting Data Engineering job:

	number	feature	weight
364	18681	etl	-3.263001
251	12488	data	-2.651070
199	10283	computer	-2.625002
748	39382	pipeline	-2.580755
375	18938	excel	-2.412889
972	52048	sql	-2.392477
845	45100	redshift	-2.355992
862	45795	report	-2.233023
89	5039	bachelor	-1.817858
463	22865	google	-1.766926

1.5.3 4.3 Decision Tree w/ TF-IDF (Model 3)

```
[102]: # train a Decision Tree Classifier
dtree_TFIDF = DecisionTreeClassifier(max_depth = 3)
dtree_TFIDF.fit(X_train, y_train)
y_pred = dtree_TFIDF.predict(X_test)
```

```
[103]: # compute y-prediction and accuracy, recall, precision, and f1 scores
y_pred = dtree_TFIDF.predict(X_test)
mod3_acc = accuracy_score(y_test, y_pred)
mod3_recall = recall_score(y_test, y_pred, average = 'macro')
mod3_precision = precision_score(y_test, y_pred, average = 'macro')
```

```

mod3_f1 = f1_score(y_test, y_pred, average = 'macro')

# print accuracy, recall, precision, f1 score, detailed report
print('-----')
print('Model 3 accuracy: ', round(mod3_acc,3))
print('Model 3 recall: ', round(mod3_recall,3))
print('Model 3 precision: ', round(mod3_precision,3))
print('Model 3 f1 : ', round(mod3_f1,3))
print('Model 3 classification report: \n',classification_report(y_test, y_pred))

```

```

-----
Model 3 accuracy:  0.648
Model 3 recall:   0.599
Model 3 precision: 0.697
Model 3 f1 :     0.611
Model 3 classification report:

```

	precision	recall	f1-score	support
0	0.79	0.58	0.67	2900
1	0.57	0.86	0.69	3025
2	0.74	0.36	0.48	1528
accuracy			0.65	7453
macro avg	0.70	0.60	0.61	7453
weighted avg	0.69	0.65	0.64	7453

```

[104]: # append model results to dataframe
model_results = model_results.append({'model': 'DecisionTree_TFIDF',
                                     'accuracy':round(mod3_acc,3),
                                     'recall': round(mod3_recall,3),
                                     'precision': round(mod3_precision,3),
                                     'f1': round(mod3_f1,3)}, ignore_index=True)

model_results

```

```

[104]:

```

	model	accuracy	precision	recall	f1
0	LogReg_OVR_TFIDF	0.872	0.863	0.853	0.857
1	LogReg_Multi_TFIDF	0.880	0.870	0.865	0.867
2	DecisionTree_TFIDF	0.648	0.697	0.599	0.611

```

[105]: fig, ax = plt.subplots(figsize=(10, 10), dpi=300)

plot_confusion_matrix(dtrees_TFIDF, X_test, y_test,
                      display_labels=['Data Scientist', 'Data Engineer',
→ 'Data Analyst'],
                      cmap=plt.cm.Blues, ax=ax)

```

```
ax.set_title('DecisionTree_TFIDF Confusion Matrix')

plt.grid(None)
plt.show()
```



1.5.4 4.4 K-Nearest Neighbors w/ TF-IDF (Model 4)

```
[106]: # setup plot to determine optimal number of neighbors
neighbors = np.arange(1, 10)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))

# Loop over different values of k
for i, k in enumerate(neighbors):
    # Setup a k-NN Classifier with k neighbors: knn
    knn = KNeighborsClassifier(n_neighbors= k)
```

```


# Fit the classifier to the training data
knn.fit(X_train, y_train)

#Compute accuracy on the training set
train_accuracy[i] = knn.score(X_train, y_train)

#Compute accuracy on the testing set
test_accuracy[i] = knn.score(X_test, y_test)

# Generate plot
plt.title('k-NN: Varying Number of Neighbors')
plt.plot(neighbors, test_accuracy, label = 'Testing Accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training Accuracy')
plt.legend()
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.show()

```



```

[107]: # train an KNN Classifier
KNN_TFIDF = KNeighborsClassifier(n_neighbors = 3)
KNN_TFIDF.fit(X_train, y_train)

```

```

[107]: KNeighborsClassifier(n_neighbors=3)

```



```
[108]: # compute y-prediction and accuracy, recall, precision, and f1 scores
y_pred = KNN_TFIDF.predict(X_test)
mod4_acc = accuracy_score(y_test, y_pred)
mod4_recall = recall_score(y_test, y_pred, average = 'macro')
mod4_precision = precision_score(y_test, y_pred, average = 'macro')
mod4_f1 = f1_score(y_test, y_pred, average = 'macro')

# print accuracy, recall, precision, f1 score, detailed report
print('-----')
print('Model 4 accuracy: ', round(mod4_acc,3))
print('Model 4 recall: ', round(mod4_recall,3))
print('Model 4 precision: ', round(mod4_precision,3))
print('Model 4 f1 : ', round(mod4_f1,3))
print('Model 4 classification report: \n',classification_report(y_test, y_pred))
```

```
-----
Model 4 accuracy:  0.716
Model 4 recall:   0.68
Model 4 precision: 0.802
Model 4 f1 :     0.697
Model 4 classification report:
              precision    recall  f1-score   support

    0           0.60       0.95       0.74       2900
    1           0.91       0.61       0.73       3025
    2           0.90       0.48       0.63       1528

 accuracy                   0.72       7453
 macro avg           0.80       0.68       0.70       7453
weighted avg           0.79       0.72       0.71       7453
```

```
[109]: # append model results to dataframe
model_results = model_results.append({'model': 'KNN_TFIDF',
                                     'accuracy':round(mod4_acc,3),
                                     'recall': round(mod4_recall,3),
                                     'precision': round(mod4_precision,3),
                                     'f1': round(mod4_f1,3)}, ignore_index=True)

model_results
```

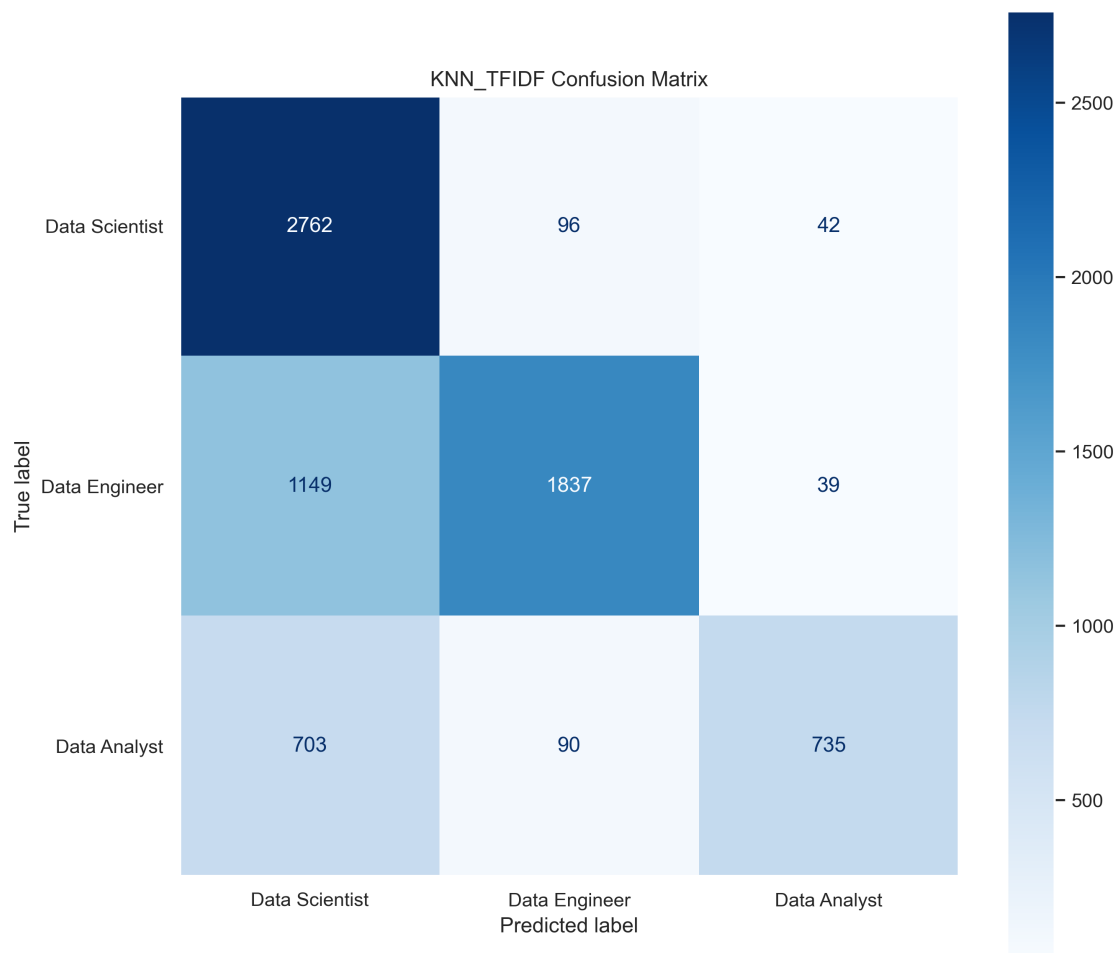
```
[109]:
```

	model	accuracy	precision	recall	f1
0	LogReg_OVR_TFIDF	0.872	0.863	0.853	0.857
1	LogReg_Multi_TFIDF	0.880	0.870	0.865	0.867
2	DecisionTree_TFIDF	0.648	0.697	0.599	0.611
3	KNN_TFIDF	0.716	0.802	0.680	0.697

```
[110]: fig, ax = plt.subplots(figsize=(10, 10), dpi=300)

plot_confusion_matrix(KNN_TFIDF, X_test, y_test,
                      display_labels=['Data Scientist', 'Data Engineer', 'Data Analyst'],
                      cmap=plt.cm.Blues, ax=ax)
ax.set_title('KNN_TFIDF Confusion Matrix')

plt.grid(None)
plt.show()
```



1.5.5 4.5 Naive Bayes w/ TF-IDF (Model 5)

```
[111]: # convert to non-sparse matrix for GaussianNB
X_train = X_train.todense()
X_test = X_test.todense()
```

```
# train a GaussianNB Classifier
GaussianNB_TFIDF = GaussianNB()
GaussianNB_TFIDF.fit(X_train, y_train)
```

```
[111]: GaussianNB()
```

```
[112]: # compute y-prediction and accuracy, recall, precision, and f1 scores
y_pred = GaussianNB_TFIDF.predict(X_test)
mod5_acc = accuracy_score(y_test, y_pred)
mod5_recall = recall_score(y_test, y_pred, average = 'macro')
mod5_precision = precision_score(y_test, y_pred, average = 'macro')
mod5_f1 = f1_score(y_test, y_pred, average = 'macro')

# print accuracy, recall, precision, f1 score, detailed report
print('-----')
print('Model 5 accuracy: ', round(mod5_acc,3))
print('Model 5 recall: ', round(mod5_recall,3))
print('Model 5 precision: ', round(mod5_precision,3))
print('Model 5 f1 : ', round(mod5_f1,3))
print('Model 5 classification report: \n',classification_report(y_test, y_pred))
```

```
-----
Model 5 accuracy:  0.613
Model 5 recall:   0.643
Model 5 precision: 0.66
Model 5 f1 :     0.611
Model 5 classification report:
              precision    recall  f1-score   support

    0           0.75         0.59         0.66         2900
    1           0.86         0.54         0.66         3025
    2           0.37         0.79         0.51         1528

 accuracy                   0.61         0.61         0.61         7453
 macro avg                 0.66         0.64         0.61         7453
weighted avg                 0.72         0.61         0.63         7453
```

```
[113]: # append model results to dataframe
model_results = model_results.append({'model': 'GaussianNB_TFIDF',
                                     'accuracy':round(mod5_acc,3),
                                     'recall': round(mod5_recall,3),
                                     'precision': round(mod5_precision,3),
                                     'f1': round(mod5_f1,3)}, ignore_index=True)

model_results
```


```
[113]:
```

	model	accuracy	precision	recall	f1
0	LogReg_OVR_TFIDF	0.872	0.863	0.853	0.857
1	LogReg_Multi_TFIDF	0.880	0.870	0.865	0.867
2	DecisionTree_TFIDF	0.648	0.697	0.599	0.611
3	KNN_TFIDF	0.716	0.802	0.680	0.697
4	GaussianNB_TFIDF	0.613	0.660	0.643	0.611

```
[114]: fig, ax = plt.subplots(figsize=(10, 10), dpi=300)

plot_confusion_matrix(GaussianNB_TFIDF, X_test, y_test,
                      display_labels=['Data Scientist', 'Data Engineer', 'Data Analyst'],
                      cmap=plt.cm.Blues, ax=ax)
ax.set_title('GaussianNB_TFIDF Confusion Matrix')

plt.grid(None)
plt.show()
```



1.5.6 4.6 Choose Best Model and Optimize Hyperparameters

```
[115]: # split train/test data 80/20
X_train, X_test, y_train, y_test = train_test_split(df_feature_TFIDF,
                                                    df_target,
                                                    train_size=0.8,
                                                    random_state=20)

# scale data
sc = StandardScaler(with_mean=False)
X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)

[116]: # choose parameters
tuned_parameters = [{'estimator__C': [100, 10, 1, 0.1, 0.01]}]

# find optimal C by grid search and fit
logreg_OVR = OneVsRestClassifier(LogisticRegression(max_iter=1000))
grid = GridSearchCV(logreg_OVR, tuned_parameters, scoring = 'f1_weighted',
                    verbose=2, cv=3)
grid.fit(X_train, y_train)

# print best score/parameter
print(grid.best_score_)
print(grid.best_params_)

# compute y-prediction
y_pred = grid.predict(X_test)
```

Fitting 3 folds for each of 5 candidates, totalling 15 fits

[CV] estimator__C=100 ...

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

/opt/anaconda3/lib/python3.8/site-

packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

/opt/anaconda3/lib/python3.8/site-

packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 47.4s remaining: 0.0s
[CV] ... estimator__C=100, total= 47.4s
[CV] estimator__C=100 ...
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
[CV] ... estimator__C=100, total= 49.1s
[CV] estimator__C=100 ...
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
[CV] ... estimator__C=100, total= 47.0s
[CV] estimator__C=10 ...
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed

```
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(

[CV] ... estimator__C=10, total= 46.9s
[CV] estimator__C=10 ...

/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
```


STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
[CV] ... estimator__C=10, total= 47.3s
[CV] estimator__C=10 ...
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/opt/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed
to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
[CV] ... estimator__C=10, total= 47.5s
[CV] estimator__C=1 ...
[CV] ... estimator__C=1, total= 38.8s
[CV] estimator__C=1 ...
[CV] ... estimator__C=1, total= 38.1s
[CV] estimator__C=1 ...
[CV] ... estimator__C=1, total= 33.7s
[CV] estimator__C=0.1 ...
[CV] ... estimator__C=0.1, total= 15.8s
[CV] estimator__C=0.1 ...
[CV] ... estimator__C=0.1, total= 14.1s
[CV] estimator__C=0.1 ...
[CV] ... estimator__C=0.1, total= 13.6s
[CV] estimator__C=0.01 ...
[CV] ... estimator__C=0.01, total= 7.4s
[CV] estimator__C=0.01 ...
[CV] ... estimator__C=0.01, total= 7.2s
[CV] estimator__C=0.01 ...
[CV] ... estimator__C=0.01, total= 7.5s

[Parallel(n_jobs=1)]: Done 15 out of 15 | elapsed: 7.7min finished
0.8764449715162215
{'estimator__C': 0.01}
```

```
[117]: mod6_acc = accuracy_score(y_test, y_pred)
mod6_recall = recall_score(y_test, y_pred, average = 'macro')
mod6_precision = precision_score(y_test, y_pred, average = 'macro')
mod6_f1 = f1_score(y_test, y_pred, average = 'macro')

# print accuracy, recall, precision, f1 score, detailed report
print('-----')
print('Model 6 accuracy: ', round(mod6_acc,3))
print('Model 6 recall: ', round(mod6_recall,3))
print('Model 6 precision: ', round(mod6_precision,3))
print('Model 6 f1 : ', round(mod6_f1,3))
print('Model 6 classification report: \n',classification_report(y_test, y_pred))
```

```
-----
Model 6 accuracy: 0.891
Model 6 recall: 0.873
```

Model 6 precision: 0.882

Model 6 f1 : 0.877

Model 6 classification report:

	precision	recall	f1-score	support
0	0.89	0.91	0.90	2900
1	0.91	0.93	0.92	3025
2	0.84	0.78	0.81	1528
accuracy			0.89	7453
macro avg	0.88	0.87	0.88	7453
weighted avg	0.89	0.89	0.89	7453

```
[118]: # append model results to dataframe
model_results = model_results.append({'model': 'LogReg_OVR_TFIDF_optimized',
                                     'accuracy': round(mod6_acc,3),
                                     'recall': round(mod6_recall,3),
                                     'precision': round(mod6_precision,3),
                                     'f1': round(mod6_f1,3)}, ignore_index=True)

model_results = model_results.sort_values(by='accuracy')
model_results
```


```
[118]:
```

	model	accuracy	precision	recall	f1
4	GaussianNB_TFIDF	0.613	0.660	0.643	0.611
2	DecisionTree_TFIDF	0.648	0.697	0.599	0.611
3	KNN_TFIDF	0.716	0.802	0.680	0.697
0	LogReg_OVR_TFIDF	0.872	0.863	0.853	0.857
1	LogReg_Multi_TFIDF	0.880	0.870	0.865	0.867
5	LogReg_OVR_TFIDF_optimized	0.891	0.882	0.873	0.877

```
[120]: fig, ax = plt.subplots(figsize=(10, 10), dpi=300)

plot_confusion_matrix(grid, X_test, y_test,
                      display_labels=['Data Scientist', 'Data Engineer', 'Data Analyst'],
                      cmap=plt.cm.Blues, ax=ax)
ax.set_title('LogReg_OVR_TFIDF_optimized Confusion Matrix')

plt.grid(None)
plt.show()
```



1.6 5. Conclusion

This project reviewed data-related job postings and tried to answer the following questions. The answers and results are shown below:

1.6.1 What tools/skills are most in demand for a *Data Analyst*?

For a Data Analyst, the top desired tools are: *sql*, *excel*, *tableau*, *python*, and *r*. The top desired skills (after parsing through some useless bigrams) are: *data analysis*, *communication skill*, and *data visualization*. See section 3.1 for the full list and plot of tools/skills frequency.

1.6.2 What tools/skills are most in demand for a *Data Engineer*?

For a Data Engineer, the top desired tools are: *sql*, *python*, *aws*, *spark*, and *azure*. The top desired skills (after parsing through some useless bigrams) are: *data pipeline*, *big data*, and *data warehouse*. See section 3.1 for the full list and plot of tools/skills frequency.

1.6.3 What tools/skills are most in demand for a *Data Scientist*?

For a Data Scientist, the top desired tools are: *python, r, sql, spark, and tableau*. The top desired skills (after parsing through some useless bigrams) are: *machine learning, data analysis, and communication skill*. See section 3.1 for the full list and plot of tools/skills frequency.

1.6.4 Which companies post the most job openings?

For Data Analyst postings, the top companies are: *ClearedJobs.Net, GEICO, CyberCoders, Booz Allen Hamilton, and Apex Systems*. See section 3.3 for the full list and plot of job company frequency.

For Data Engineer postings, the top companies are: *Amazon, Optello, Facebook, Apple, and CyberCoders*. See section 3.3 for the full list and plot of job company frequency.

For Data Scientist postings, the top companies are: *Amazon, Facebook, Booz Allen Hamilton, Apple, and Optello*. See section 3.3 for the full list and plot of job company frequency.

1.6.5 Can a classifier be built which predicts job role/title (Data Analyst, Data Scientist, or Data Engineer) based on job description?

Five different initial models were chosen for the multi-class (three class) classification problem. These models included the following: Logistic Regression (One vs. All), Logistic Regression (Multinomial), K-Nearest Neighbor, Decision Tree, and Naive Bayes. Based on an initial fitting of the model types to the classification problem and collecting model metrics, the Logistic Regression model performed the best. This model type was then further optimized via a hyperparameter grid search. The resulting best model, LogReg_OVR_TFIDF_optimized, is shown below with model performance metrics

Model	Accuracy	Precision	Recall	f1
GaussianNB_TFIDF	0.613	0.660	0.643	0.611
DecisionTree_TFIDF	0.648	0.697	0.599	0.611
KNN_TFIDF	0.716	0.802	0.680	0.697
LogReg_OVR_TFIDF	0.872	0.863	0.853	0.857
LogReg_Multi_TFIDF	0.880	0.870	0.865	0.867
LogReg_OVR_TFIDF_optimized	0.891	0.882	0.873	0.877

Next steps for improving model performance include trying additional model types, additional hyperparameter tuning (e.g. more solvers, more C estimators, more penalty terms, etc), review imbalanced classes and upsample/downsample appropriately, and try different NLP cleaning approaches (e.g. add more stopwords).